

# Efficient Duplicate Elimination in SPARQL to SQL Translation

Dimitris Bilidas and Manolis Koubarakis

Department of Informatics and Telecommunications  
National and Kapodistrian University of Athens, Greece  
`{d.bilidas,koubarak}@di.uoa.gr`

**Abstract.** Redundant processing is a key problem in SPARQL to SQL query translation in *Ontology Based Data Access* (OBDA) systems. Many optimizations that aim to minimize this problem have been proposed and implemented. However, a large number of redundant duplicate answers are still generated in many practical settings, and this is a factor that impacts query execution. In this work we identify specific query traits that lead to duplicate introduction and we track down the mappings that are responsible for this behavior. Through experimental evaluation using the OBDA system Ontop, we exhibit the benefits of early duplicate elimination and show how to incorporate this technique into query optimization decisions.

## 1 Introduction

*Ontology Based Data Access* (OBDA) is a database technique in which an ontology is linked to underlying data sources through mappings. An end user can pose queries over the ontology, which we assume to represent a familiar vocabulary and conceptualization of the user domain. The OBDA system automatically translates the query and sends it for execution to the underlying data sources, providing the end user with a convenient abstraction over possibly complex schemas and details about the data storage and query processing. The query translation involves *query rewriting*, where the initial query is transformed in order to take into consideration the ontology axioms, and *query unfolding* where the rewritten query is transformed into another query expressed in the query language of the underlying data sources. In what follows we consider an OBDA setting, where an OWL ontology is linked through mappings to data stored in a *relational database management system* (RDBMS) providing the user with access to a virtual RDF graph. The original query is expressed over this virtual RDF graph in the SPARQL query language, and the result of rewriting and unfolding is a SQL query. As an example consider the relational tables from Figures 1a and 1b and the mappings from Figure 1e. In these mappings *hasDirector* and *hasActor* are properties defined in the ontology, whereas *f*, *g* and *h* are functions roughly corresponding to string concatenation. These functions are responsible for constructing an object that acts as an ontology instance out of values occurring in the database. In our setting, they construct an RDF term represented by an IRI.

If an RDF graph is lean (i.e., if it has no instance which is a proper subgraph of itself), evaluation of basic graph pattern on this graph can never result in duplicate answers [16]. However, duplicates in SPARQL query evaluation can be introduced from the operations of projection and union. Therefore, it is crucial that SQL query fragments resulting from the translation of basic graph patterns are also duplicate free when evaluated in a source database. Note also that, as there is no restriction regarding the multiplicity of the results of SQL queries used in a mapping, duplicates may be introduced even during evaluation of a single triple pattern.

In this setting, duplicates are redundant answers whose impact can be detrimental for query evaluation, as the size of intermediate results can increase exponentially in the number of input tables, in a query that involves several joins. Even if the final SQL query produced by an OBDA system dictates that the result should be duplicate free using the SQL `DISTINCT` or `UNION` keyword, relational systems rarely consider early duplicate elimination in order to limit the size of intermediate results, but only perform the task on the final query result. This behavior is justified by the fact that duplicate elimination is a costly blocking operation [2] and also that the SQL queries are usually formulated by expert users who take into consideration the integrity constraints of normalized relational schemas. Under these assumptions, considering early duplicate elimination options during optimization is not usually regarded worthy. Contrary to this situation for SQL queries, it has been ascertained [10] that in real world OBDA settings, duplicate answers frequently dominate query results and also that this appears as “noise” to end users that might be using a visual query formulation tool. In what follows we briefly identify reasons for this behavior.

In a normalized relational schema duplicates usually show up in SQL queries due to projections. As long as an OBDA mapping that generates virtual triples uses a column or combination of columns whose values are unique, e.g. a primary key of a table as the subject or object of each triple, then all triples coming from this mapping will be distinct. However, a denormalized schema or mappings that do not use unique columns can lead to duplicate introduction even from a single triple pattern. One more case in which duplicates are introduced has to do with the use of property *rdf:type*. Mappings defining *rdf:type* property should also be duplicate free, but as OBDA systems employ reasoning with respect to domain and range of properties, part of the mapping that defines a property could be used for evaluating a triple pattern with *rdf:type*. Then the other part of the mapping is projected out, something that can lead to a large number of duplicates. This is illustrated in Example 1.

*Example 1.* Consider the mappings from Figure 1e and also that the ontology defines the range of property *hasDirector* to be the class *Director*. Now consider a triple pattern *?x rdf:type Director* that asks for the entities that belong to the class *Director*. One way to obtain such entities is to use the entities participating as objects to the property *hasDirector*. As we are not interested in the subjects of the generated triples, we are projecting out the *title* column, adding only a condition that this column should not be null. The resulting SQL to be sent for

evaluation is shown in Figure 1c (where  $||$  represents the string concatenation operator in SQL)

In the example above, the number of duplicate results is equal to the average number of movies that correspond to each director in the *movies* table. This can become a heavy burden for query evaluation, especially when the query involves several joins, as all these duplicate values need to be joined with other tables, that may as well contain redundant values, and as a result, the overall cost can be increased by orders of magnitude, depending on the exact number of duplicates. A second issue is that, as the final answers need to be distinct, a duplicate elimination has to be performed on the final query result that consists of tuples of RDF terms instead of database values, making the relational engine perform this over (usually large) string values. As a typical query translation produces a UNION SQL query, results from different subqueries have to be merged and deduplicated.

In this work we present efficient solutions to the described problems, considering ontologies belonging to the OWL 2 QL language<sup>1</sup>, as the W3C recommendation for query answering against datasets stored in relational back-ends. Nevertheless, several aspects of this work can be considered for other ontology languages as well. We start by providing some preliminaries regarding ontologies, mappings and relational databases (Section 2), including a condition that guarantees distinct results for specific queries, by taking into consideration functional database dependencies. We then proceed to describe a process, that given an initial mapping collection, identifies the (possibly modified) mapping assertions which are responsible for introduction of duplicates (Section 3). Deduplicated results of these assertions can be materialized offline and replace the original assertions during query execution. In case materialized views are not a viable option, for example because of read-only access to data or for efficiency reasons, we proceed to deal with the problem of duplicates during query execution.

In Section 4, we describe how a structural optimization of pushing duplicate elimination before IRI construction can be applied and obtain an equivalent query with the one produced during unfolding. Then, in Section 5 we propose a heuristic that can help an OBDA system acting outside the relational engine to take decisions regarding the duplicate elimination of intermediate results that have been produced by the previously identified mapping assertions. This heuristic uses only information regarding the size of source relations and an estimation of the final size of the query, using easily obtained data summarization information from the underlying relational engine. In Section 6 we propose an improvement to self-join elimination methods currently employed by OBDA systems, relevant to redundancy due to duplicates, as it is applied in cases of denormalized relational schemas, or mappings that do not use primary keys, which both lead to introduction of duplicates.

Having implemented our optimizations in the state of the art OBDA system Ontop [19], in Section 7 we present their impact on NPD and LUBM benchmarks

<sup>1</sup> <https://www.w3.org/TR/owl2-profiles/>

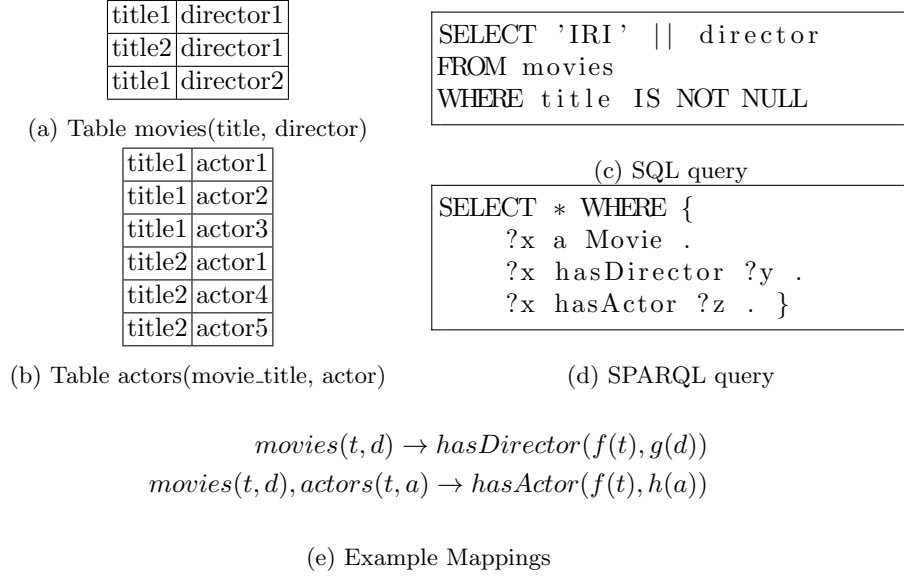


Fig. 1: Example Database

using three different relational back-ends. Specifically we show that duplicates are present in many mappings from both benchmarks and that specific queries are heavily impacted by this fact, leading to evaluation times of more than 1000 seconds, which in some cases can be reduced to a few seconds. Also we show that, excluding those queries that give a timeout of more than 1000 seconds in the unoptimized setting, the overall improvement for the rest of the queries can be more than 60%. We also show that our self-join optimization affects several queries from both benchmarks and that pushing duplicate elimination before IRI construction, apart from reducing the overall execution time, has also the effect of obtaining the answers for many queries in a pipelined fashion, leading specifically to large improvement in the time of obtaining the first answers. Finally, we evaluate our heuristic and show that its usage is justified and that for query mixes from the two benchmarks, such that low selectivity queries do not dominate execution time, it can lead to overall improvement of up to 25% compared to the strategy of always performing duplicate elimination. In Section 8 we present relevant work and conclusions.

## 2 Preliminaries

We consider the following pairwise disjoint alphabets:  $\Sigma_O$  of ontology predicates,  $\Sigma_R$  of database relation predicates,  $Const$  of constants,  $Var$  of variables and  $A$  of function symbols where each function symbol has an associated arity. We also consider that  $Const$  is partitioned into  $DB_{Const}$  of database constants and  $\mathcal{O}_{Const}$  of ontology constants.

As in [18], we use functions with symbols from  $\mathcal{A}$  in order to solve the so called *impedance mismatch problem* of constructing ontology objects from values occurring in the database. These functions roughly correspond to IRI templates specified in the R2RML<sup>2</sup> language.

## 2.1 Databases.

We start by giving definitions for database instances and queries over them, following the bag semantics from [4].

A *bag*  $B$  is a pair  $(A, \mu)$ , where  $A$  is a set called the underlying set of  $B$  and  $\mu$  is a function from elements of  $A$  to the positive integers, which gives the multiplicities of elements of  $A$  in  $B$ . The *underlying set* of a bag  $B$  will be denoted by  $US_B$ .

A *relation instance* is a bag of tuples of fixed arity using constants from  $DB_{Const}$ .

A *source schema*  $S$  is a set of relation names from  $\Sigma_R$ .

A *database instance*  $D$  for a source schema  $S$  is a mapping from relation names in  $S$  to relation instances.

## 2.2 Queries.

We define queries following the bag semantics of [4]. In our definitions we use the term “SQL query” although the syntax of our formulas is that of first-order logic. Similarly, relation instances are viewed as bags of *ground atoms* (i.e., with no variables) of first-order logic.

A *SQL query* over a relational schema  $S$  is an expression that has the form:  $Query(\mathbf{x}) \leftarrow \alpha$ , where  $\alpha$  is a first order expression containing predicates from  $\Sigma_R$ , which are among the relations that belong to  $S$ ,  $Query \in \Sigma_R$ ,  $Query \notin S$  and  $\mathbf{x}$  is a vector of constants from  $DB_{Const}$  and variables from  $Var$  that appear in  $\alpha$ .

A *conjunctive query*  $CQ$  over a relational schema  $S$  is a SQL query, where  $\alpha$  has the form  $R_1(\mathbf{x}_1) \wedge \dots \wedge R_n(\mathbf{x}_n)$ , where  $\mathbf{x}_1, \dots, \mathbf{x}_n$  are vectors of constants from  $DB_{Const}$  and variables from  $Var$ , and  $R_1, \dots, R_n \in S$ . Variables from  $\mathbf{x}_1, \dots, \mathbf{x}_n$  that do not appear in  $\mathbf{x}$  are existentially quantified, but we omit the quantifiers in order to simplify the reading. CQs roughly correspond to SQL Select-From-Where queries.

Let  $q$  be the SQL query  $Query(\mathbf{x}) \leftarrow \alpha$ , we will denote by  $\prod_i(q)$  the query resulting from the projection of the  $i$ -th answer term of  $q$ , that is the query:  $Query(x_i) \leftarrow \alpha$

An *assignment mapping* of a conjunctive query  $Q$  into a database instance  $D$  is an assignment of values from  $DB_{Const}$  belonging to  $D$  to the variables of  $Q$  such that every atom in the body of  $Q$  is mapped to a ground atom in  $D$ . Let  $\theta$  be an assignment mapping of  $Q$  into database instance  $D$  and let  $X$  be a

<sup>2</sup> <https://www.w3.org/TR/r2rml/>

variable in  $Q$ . We denote by  $\theta(X)$  the constant in  $DB_{Const}$  to which  $\theta$  maps  $X$  and we denote by  $\theta(R_i(\mathbf{x}_i))$  the ground atom to which  $R_i(\mathbf{x}_i)$  is mapped.

Let  $\mu_i$  denote the multiplicities  $\mu(\theta(R_i(\mathbf{x}_i)))$ ,  $i = 1, \dots, n$ . The *result* due to  $\theta$  of a conjunctive query  $Q$  over  $D$  is the atom  $(\theta(\mathbf{x}), m_Q)$  with the multiplicity  $\mu_Q = \mu_1 \mu_2 \cdots \mu_n$ .

The *result* of a conjunctive query  $Q$  over a database instance  $D$  denoted by  $Q(D)$  is given by  $\cup_{\theta} r_{\theta}$ , where  $\theta$  is any assignment mapping of  $Q$  into  $D$  and  $r_{\theta}$  is the result due to  $\theta$ .

### 2.3 Ontology and Mappings.

A *TBox* is a finite set of ontology axioms.

An *ABox* is a finite set of membership assertions  $A(\rho)$  or role filling assertions  $P(\rho, \rho')$ , where  $\rho, \rho' \in \mathcal{O}_{Const}$  and  $A, P \in \Sigma_O$  denote a concept name and role name respectively.

A DL ontology  $\mathcal{O}$  is a pair  $\langle \mathcal{T}, \mathcal{A} \rangle$  where  $\mathcal{T}$  is a *TBox* and  $\mathcal{A}$  an *ABox*.

A *mapping assertion*  $m$  from a source schema  $S$  to a *TBox*  $\mathcal{T}$  has the form:  $\phi(\mathbf{x}) \rightarrow \psi_1 \wedge \dots \wedge \psi_n$ , where  $\phi(\mathbf{x})$  will be denoted by  $body(m)$  and it is a SQL query over a database instance  $D$ , each  $\psi_i$  has the form  $P_i(cc_i^1(\mathbf{x}_i^1), cc_i^2(\mathbf{x}_i^2))$  or  $C_i(cc_i^1(\mathbf{x}_i^1))$  with  $P_i$  ( respectively  $C_i$ )  $\in \Sigma_O$  a property (respectively concept) name, and each  $cc_i^j \in A$  is a function with arity equal to the length of  $\mathbf{x}_i^j$  and range a subset of  $\mathcal{O}_{Const}$ . The conjunction in the right hand side will be denoted by  $head(m)$ . A *mapping*  $\mathcal{M}$  is a finite set of such mapping assertions.

In this paper we consider *global-as-view* (GAV) mappings, where all variables in  $\psi_1 \wedge \dots \wedge \psi_n$  also appear in  $\mathbf{x}$ . In this setting,  $\mathcal{M}$  can be transformed into a set of equivalent mapping assertions where the head of each assertion consists of a single atom [18]. In what follows we assume that the head of every mapping assertion consists of a single atom.

Let  $\mathcal{M}$  be a mapping, we will use the symbol  $\mathcal{M}_{CQ}$  to denote the assertions from  $\mathcal{M}$  whose body is a CQ over the database schema.

In correspondence with CQs over a relational schema, we define a CQ over an ontology  $\mathcal{O}$  as an expression of the form:  $Query(\mathbf{x}) \leftarrow P_1(\mathbf{x}_1) \wedge \dots \wedge P_n(\mathbf{x}_n)$  where  $\mathbf{x}_1, \dots, \mathbf{x}_n$  are vectors of constants from  $\mathcal{O}_{Const}$  and variables from  $Var$ ,  $\mathbf{x}$  is a vector of constants from  $\mathcal{O}_{Const}$  and variables from  $Var$  that appear in  $\mathbf{x}_1, \dots, \mathbf{x}_n$ , and  $P_1, \dots, P_n \in \Sigma_O$  are ontology predicates that appear in  $\mathcal{O}$ .

A *union of conjunctive queries*  $UCQ$  over an ontology  $\mathcal{O}$  is an expression of the form  $Query(\mathbf{x}) \leftarrow CQ_1(\mathbf{x}) \vee \dots \vee CQ_n(\mathbf{x})$ , where each  $CQ_i$  for  $i = 1 \dots n$  is an expression of the form  $P_1^i(\mathbf{x}_1^i) \wedge \dots \wedge P_n^i(\mathbf{x}_n^i)$  as in the previous definition.

### 2.4 Primary Keys and Duplicate Free Results.

The *duplicate-tuple ratio*  $DTR_R$  of a relation instance  $R$  is equal to  $\frac{\sum_{t \in US_R} \mu(t)}{|US_R|}$ . A relation instance with  $DTR$  equal to 1, will be called a *duplicate-free* relation instance.

As a first step, we want to take advantage of primary key constraints that hold in the source schema, so that we will avoid estimating DTR for relations which can be deduced to be duplicate-free. It is known that the interaction of key dependencies and functional dependencies is different under set and bag semantics [12]. In what follows we are only interested in key dependencies, since only these can be proved useful in order to decide if a result of a conjunctive query is duplicate-free, and also they correspond to the primary key declarations of SQL.

Let  $R$  be a relation of arity  $k$  in a source schema  $S$  and  $X, Y$  subsets of  $\{1, 2, \dots, k\}$  with  $X \cap Y \subseteq \emptyset$  and  $X \cup Y = \{1, 2, \dots, k\}$ . For an atom  $R(\mathbf{a})$  the *projection* of terms using  $X$  is the set of terms from  $\mathbf{a}$  whose subscripts correspond to the elements of  $X$ . We will denote this set as  $\pi_X(R(\mathbf{a}))$ . We say that a *key dependency*  $X \rightarrow Y$  holds for  $R$  in a database instance  $D$  of  $S$ , if there are no distinct atoms  $R(\mathbf{a})$  and  $R(\mathbf{b})$  in  $D$  such that  $\pi_X(R(\mathbf{a})) = \pi_X(R(\mathbf{b}))$  and for every atom  $R(\mathbf{a})$  in  $D$ ,  $\mu(R(\mathbf{a})) = 1$ . It is straightforward to see that any instance of a relation for which there is a primary key, is duplicate-free.

Later on, in Section 3, when we will want to see if the result of a mapping assertion is duplicate-free, we will use the following important proposition of [15]:

**Proposition 1.** *Let  $Q$  be a conjunctive query over a source schema  $S$ . Then, if for each  $R_i(\mathbf{u}_i)$  in the body of  $Q$  there is a key dependency  $X \rightarrow Y$  in  $S$  and for each element  $u_{i_j} \in \pi_X(R(\mathbf{u}_i))$ , either  $u_{i_j} \in DB_{Const}$  or  $u_{i_j}$  appears in  $Query(\mathbf{x})$ , then the result of  $Query(\mathbf{x})$  is a duplicate-free relation instance on any database instance for  $S$ .*

## 2.5 Query Rewriting and Unfolding

As we mentioned in Section 1, query answering in OBDA consists of query rewriting and query unfolding. During query rewriting, an initial CQ over an ontology is rewritten in order to take into consideration the ontological axioms. The result of this process is a query, that when posed over the ABox (that is by disregarding all the ontological axioms), will return the same answers as the initial query posed over the ontology. This is done using the notion of *certain answers*, that is answers present in every model of the ontology [18]. We omit details, as in this work we mainly consider the result query of this process. We just need to note that several methods exist for query rewriting over OWL 2 QL ontologies. In this work we consider that the result of the query rewriting step is a UCQ over the ontology, as it is the case for several rewriting methods [18, 11, 5, 17]. We discuss applicability on rewriting methods producing different forms of queries in Section 8.

Regarding query unfolding with respect to a Mapping  $\mathcal{M}$ , a method based on partial Datalog evaluation with functional terms is presented in [18]. As before, we omit a detailed description and note that the result of this process is a query over the relational schema that has the form

$$Query(\mathbf{x}) \leftarrow Q_1(\mathbf{x}) \vee \dots \vee Q_n(\mathbf{x}) \quad (1)$$

where each  $Q_i$  for  $i = 1 \dots n$  is an expression of the form

$$Q_i(\mathbf{f}_i(\mathbf{x}_i)) \leftarrow Aux_1(\mathbf{x}_1^i) \wedge \dots \wedge Aux_l(\mathbf{x}_l^i)$$

where each  $f_i^j \in \mathbf{f}_i$  is a function whose function name belongs in  $\Lambda$  and whose variable arguments are among the variables of  $\mathbf{x}_1^i, \dots, \mathbf{x}_l^i$  and each  $Aux_j$  for  $j = 1 \dots l$  corresponds to  $body(m)$  for some  $m \in \mathcal{M}$ .

### 3 Offline Duplicate Elimination With Materialized Views

One solution to the problem of duplicates, is to track down the mapping assertions which are responsible for duplicates, create materialized views with the distinct results, possibly with indexes, and then use these views instead of the original assertions during query unfolding. It is reasonable to expect that this solution will give the best performance during query execution, but on the other hand this incurs expensive preprocessing and also, using materialized views in the database increases the database maintenance load, especially for frequently updated tables, as well as the the database size. Also, this solution will not take into consideration duplicates due to projections in the SPARQL query and finally, it is not in line with the overall approach of providing the end user with access to several underlying data sources, without the need to modify data, and on a practical level, such access may not be even possible.

Nevertheless, even in the case where one chooses to use materialized views, it is not straightforward exactly which of the mapping assertions should be chosen. In the rest of this section we describe a process to find the exact assertions for this setting, whereas in the following sections we consider the case where no materialization happens and all processing needs to be done during query execution.

Given a mapping  $\mathcal{M}$  and a database instance  $D$  over a schema  $S$ , a straightforward solution is to materialize all  $m \in \mathcal{M}$  such that  $DTR_{head(m)}(D) > 1$ , but as the query produced after rewriting takes into consideration the ontology axioms, implied assertions may be used, such that a specific variable has been projected out from the outputs of the body of an existing assertion due to reasoning for class instances with respect to domain or range of a property like in Example 1, where the modified mapping implied assertion is the following:

$$movies(t, d) \rightarrow Director(d)$$

The exact way that this implied assertion will be used depends on the rewriting method, but in any case, in the resulted SQL query column *title* will not be in the *Select* clause.

Situations like this can be identified offline, by analyzing the ontology and the original mapping. The first step is to find the ontology axioms of the form  $\exists P.\top \sqsubseteq C$  and  $\exists P^-\top \sqsubseteq C$  (or equivalently  $\top \sqsubseteq \forall P.C$ ) that define the domain and range of some property  $P$  to be a class  $C$  in our ontology. Then we identify the mapping assertion in  $\mathcal{M}$  that generates RDF triples which have as predicate the property  $P$  and we modify the target SQL query of the mapping, by projecting out the columns used for the subject or object respectively. At this point, we can skip certain mappings that are covered by the corresponding *rdf:type*



mapping for a given class, as OBDA systems eliminate the usage of the property triple pattern for these cases based on foreign key relationships [19]. Consider again the case of Example 1: if we had one more database table *director\_info* which has a primary key *id* and we also know that *director* in *movies* is a foreign key that references this primary key, then if we also had the mapping:

$$director\_info(id, \dots) \rightarrow Director(f(id))$$

the previously obtained mapping can be skipped as it is redundant and will not be used by the OBDA system.

The method is described in Algorithm 1. *ComputeDTR* is a function that returns the DTR for the query passed as argument. If  $DTR = 1$  according to proposition 1, then access to data is avoided altogether, otherwise the actual DTR is computed by sending two count queries: with and without the distinct modifier. Later, when the DTR needs to be determined during query optimization, an estimation based on data summarization is used instead (Section 5). Function *ExistsFK* returns true if a foreign key exists between the output column of query *query*<sub>1</sub> and the output column of query *query*<sub>2</sub>. The result of this algorithm is a set of mapping assertions, possibly annotated with information about the projection of a column. Modification of the produced SQL query in order to take into consideration the views created for these mappings, instead of the original body of the mapping, can simply be performed in the final step of the query unfolding, where each *Aux*<sub>*j*</sub> is replaced by the corresponding SQL, and as a result it is independent of the query rewriting method.

## 4 Pushing Duplicate Elimination Before IRI Construction

In SPARQL to SQL approaches, pushing joins inside unions is a well known *structural* optimization, so that joins over IRIs are avoided and relational columns, whose values are possibly indexed, are used instead. Methods for unfolding based in partial datalog evaluation (see Section 2.5) produce such queries, where additionally, union subqueries that contain joins between incompatible IRIs, that when evaluated will produce an empty result, are completely discarded. Also, the *safe separator*<sup>3</sup> of the R2RML mapping language can be used to ensure that concatenation of multiple columns cannot produce the same value with that of a single column [20]. In a similar manner, it can be very useful to perform duplicate elimination before IRI construction. In this section we discuss the process of transforming the unfolded query that has the form shown in formula (1) at page 7, into an equivalent one, such that duplicate elimination is performed on database values.

To do so, we must group together union subqueries that have the same select clause up to variable (column name) renaming. In case of groups that contain only one subquery, Proposition 1 can be used to avoid duplicate elimination altogether. In our case the situation is more complicated, as we want to ensure that tuples produced from different IRI templates cannot possibly have equal values. Consider for example the following query:

<sup>3</sup> <https://www.w3.org/TR/r2rml/#dfn-safe-separator>

**Algorithm 1:** Track down SQL queries that contain duplicates

---

```

1 MappingsWithDuplicates (mapping  $\mathcal{M}$ , ontology  $\mathcal{O}$ , database schema  $S$ ,
   database instance  $D$ );
   Output: Mapping assertions possibly annotated with a projected column
   Uses :  $ComputeDTR(query, db\_schema, db\_instance)$ 
   Uses :  $ExistsFK(schema, query_1, query_2)$ 
2  $result := \emptyset$ ;
3 for  $m \in \mathcal{M}$  do
4   if  $head(m)$  is Class assertion then
5     if  $ComputeDTR(body(m), S, D) > 1$  then
6       | add  $m$  to  $result$ ;
7     end
8   else
9     /*  $head(m)$  is Property assertion with predicate  $P$  */
10    if  $\mathcal{O}$  contains  $\exists P.\top \sqsubseteq C$  and
11       $ComputeDTR(\prod_1(body(m)), S, D) > 1$  and not  $(\exists m2 \in \mathcal{M} \text{ s.t.}$ 
12         $predicate\ of\ head(m2)\ is\ C\ and$ 
13         $ExistsFK(S, \prod_1(body(m)), body(m2)))$  then
14        | add  $m$  to  $result$  for projection of 1st column;
15      end
16    if  $\mathcal{O}$  contains  $\exists P^-\top \sqsubseteq C$  and
17       $ComputeDTR(\prod_2(body(m)), S, D) > 1$  and not  $(\exists m2 \in \mathcal{M} \text{ s.t.}$ 
18         $predicate\ of\ head(m2)\ is\ C\ and$ 
19         $ExistsFK(S, \prod_2(body(m)), body(m2)))$  then
20        | add  $m$  to  $result$  for projection of 2nd column;
21      end
22    end
23 end
24 return  $result$ ;

```

---

```

SELECT ':Person' || alias1.id AS x
FROM table1 alias1
UNION
SELECT ':Person' || alias1.key AS x
FROM table2 alias1
UNION
SELECT ':Person' || alias1.id ||
      '/' || alias1.name AS x
FROM table3 alias1

```

Given that the '/' character is a safe separator, the third subquery cannot produce any result tuple that will be the same with a result tuple coming from the first two subqueries. On the other hand, there is a possibility that the first two

subqueries may produce the same answer. The following rewriting of this query can be used:

```

SELECT  ':Person' || var1 AS x,
FROM
  (SELECT
    alias1.id as var1
  FROM
    table1 alias1
  UNION
  SELECT
    alias1.key AS var1
  FROM
    table2 alias1
  )
UNION ALL
SELECT DISTINCT  ':Person' ||
  alias1.id || '/' || alias1.name AS x,
FROM table3 alias1

```

Note that UNION ALL operator is simply concatenating the results. When the UNION ALL is the outer operator of a query, it is reasonable for the RDBMS to start sending the results in a pipelining fashion, as they are produced from each subquery without saving or waiting for all the results to be produced. In this sense, it can be considered a “cheap” operator in contrast to UNION. If we know that column *id* of *table3* is a primary key, then the distinct keyword of the last subquery can be eliminated. Even if this is not the case, the resulted query has several advantages over the initial. First, the duplicate elimination process has been separated over two distinct result sets and also each tuple is smaller in size. This gives to the RDBMS the opportunity to better utilize available memory, as it now has smaller datasets to perform duplicate elimination, or even parallelize the process. Available indexes on the columns can be used. Also, as discussed, when there is no blocking outer operator, results are produced in a pipelined fashion. This way the first results can be obtained very quickly and, as IRI construction is an expensive operation, the difference can be impressive when we have large results and the processing for each subquery is relatively cheap.

Following the described process, starting from a query as in formula 1, we obtain a query that has the form

$$Query(\mathbf{x}) \leftarrow UCQ_1(\mathbf{x}) \vee \dots \vee UCQ_l(\mathbf{x}) \quad (2)$$

where each  $UCQ_i$  for  $i = 1 \dots l$  is an expression of the form

$$UCQ_i(\mathbf{f}_i(\mathbf{v}_i)) \leftarrow Q_i^1(\mathbf{v}_i) \vee \dots \vee Q_i^l(\mathbf{v}_i) \quad (3)$$

where  $\mathbf{v}_i$  is a vector of variables not occurring in (1) and  $Q_i^1(\mathbf{v}_i) \vee \dots \vee Q_i^l(\mathbf{v}_i)$  result from exactly those conjuncts of (1) that have  $\mathbf{f}_i$  in the left hand side,

by replacing each variable in  $\mathbf{x}_i$  with the variable in the same position in  $\mathbf{v}_i$  and adding a conjunction of variable equalities  $EQ$  between variables of  $\mathbf{v}_i$  that correspond to positions of  $\mathbf{x}_i$  where the same variable occurred. That is, each  $Q_i^j$  for  $j = 1 \dots k$  has the form:

$$Q_i^j(\mathbf{v}_i) \leftarrow Aux_1(\mathbf{v}_i^1) \wedge \dots \wedge Aux_n(\mathbf{v}_i^k) \wedge EQ. \quad (4)$$

In the corresponding SQL query, disjunctions in (2) can be translated to UNION ALL and only disjunctions in (3) need to be translated to UNION (or add DISTINCT if there is only one disjunct) avoiding duplicate elimination over IRIs. Also,  $EQ$  can be replaced by choosing one variable for each equality and replacing all occurrences with the specific variable, and then add a renaming operator on the *Select* clause corresponding to  $Q_i^j(\mathbf{v}_i)$ .

A problem that may arise is that it is not always possible to perform this structural optimization, because queries with different select clauses can still produce the same result. Suppose that the second subquery in the previous example had the following select clause:

```
SELECT ':Person1' || alias1.id AS x.
```

As `:Person` is a substring of `:Person1`, the second subquery may still produce the same result as the first one, but in this case we cannot just perform the duplicate elimination on the *id* columns of the two tables. Such cases are peculiar in the sense that a substring relation must hold between constant arguments of different IRI templates. In order to identify exactly when two IRI templates can produce the same result, one should construct a regular expression from each template and then examine if the intersection of the languages produced by these regular expressions is empty. Given that for most cases it is straightforward to make sure that two different IRI templates cannot generate the same result, we instead use a set of quick checks, such that if any of these checks succeed, then it is guaranteed that the IRI templates cannot generate the same result. These include a check for inequality in the number of safe separator characters, a check that the first or last arguments of two templates are not substrings etc. If none of these checks succeed, then we cannot be sure if these IRI templates can produce the same result or not, and we do not push duplicate elimination before IRI construction for the specific subqueries only.

## 5 Incorporating Decisions for Duplicate Elimination in Query Execution

In this section we consider a query that has the form shown in formula (2) and describe a process that decides about early duplicate elimination. We start by considering separately each union subquery that has the form shown in formula (4). If we consider  $Aux_1, \dots, Aux_n$  to be relation names belonging to the database schema, this is a CQ, but in practice  $Aux_1(\mathbf{v}_1^i), \dots, Aux_n(\mathbf{v}_k^1)$  are arbitrary FOL queries. Our method relies on an estimation about the final result size of each union subquery. To obtain this estimation, DTR for every mapping assertion in

$\mathcal{M}$  is no longer enough, as it was the case when creating materialized views, but we should gather some statistics from the database in the form of data summarization for all the columns that can be possibly referenced from a query, that is all the columns in the SQL query of some mapping assertion. As making an estimation for an arbitrary FOL query is an involved process, we make a distinction between assertions in  $\mathcal{M}_{CQ}$  and assertions in  $\mathcal{M} \setminus \mathcal{M}_{CQ}$ . We consider that the latter are primitive tables as if they were virtual views, and we collect statistics only for the output columns, whereas the former are parsed and we collect statistics for all the referenced columns.

As a result, we can consider that each union subquery has the form

$$Q(\mathbf{x}) \leftarrow Aux_1(\mathbf{x}_1) \wedge \dots \wedge Aux_n(\mathbf{x}_n) \wedge R_1(\mathbf{x}_{n+1}) \wedge \dots \wedge R_m(\mathbf{x}_{n+m}) \quad (5)$$

where each  $Aux_1, \dots, Aux_n$  corresponds to  $body(m)$  for some mapping assertion  $m \in \mathcal{M} \setminus \mathcal{M}_{CQ}$  and  $R_1, \dots, R_m$  are relation names from the database schema. We will refer to each conjunct in the right hand side of (5) as an *input table* of query  $Q(\mathbf{x})$ .

Let  $q$  be a query as in (5) and  $I_i(\mathbf{x}_i)$  be an input table of  $q$ . The query  $ans(\mathbf{x}_{c_i}) \leftarrow I_i(\mathbf{x}_i)$ , where  $\mathbf{x}_{c_i}$  contains exactly the variables of  $\mathbf{x}_i$  that appear more than once in  $q$ , will be called the *projection query* of input table  $I_i(\mathbf{x}_i)$  from  $q$ . Additionally, let  $D$  be a database instance (which will be implied).

*Analyzing External Tables.* As we operate outside the RDBMS engine, in order to extract the needed information we should import all the corresponding data, something that is clearly not practical. Luckily we have several other options. One first option is to only import a random sample and extract the needed information from that, as most database vendors support ordering the results by a random function. One second option is to obtain the data summarization directly from the RDBMS, if it provides a way to access this information. This option is likely to give the most accurate results, but it is highly depended on the specificities of each database vendor. One third option is to build a simple single-bucket histogram for each column, by sending for execution queries that ask for the number of values, number of distinct values, minimum and maximum value. Simple histograms like this are known to give imprecise selectivity estimations for filter and join results of attributes that exhibit skewness [8], but on the other hand their construction and usage is faster in comparison to more elaborate kinds of histograms. For our experiments we have chosen the last option, as it is fast and simple and can be applied to any underlying RDBMS. This is an offline process, that needs to be done before query execution, similar to an analyze command in a database schema, as it only depends on mappings and data. Adopting the commonly used value independence assumption between the result attributes and the uniformity of values in an attribute [24], we estimate the distinct tuples of the relation to be the product of the distinct values of its attributes. In case this value is larger than the number of tuples in the relation, we estimate that all tuples are distinct.

### 5.1 Early Duplicate Elimination of Intermediate Results

Let us suppose that we have a single SQL subquery coming from the translation of a SPARQL query and we have to take the decision regarding a single input table (either “real” primitive table or virtual view) used in this subquery; we will take into consideration different union subqueries in Section 5.2. In this case, it may be profitable to dictate the RDBMS to perform the duplicate elimination on projection query of the specific input table at the beginning of query execution, store the duplicate-free intermediate result in a temporary table and use it for the specific query. This can be done in several ways depending on the exact SQL dialect and capabilities of the underlying system. For example one can use (non-recursive) common table expressions or temporary table definitions. Of course the exact decisions as to when this should happen depend on several factors, including the exact query, the DTR of the projection query of the input table, the number of uses of the specific input table in the query, the choice to save the temporary table in disk or keep it in memory and several other factors that depend on the database physical design, database tuning parameters, the exact query execution plan and the evaluation methods chosen by the optimizer of the RDBMS. Uncertainty about query execution costs is an inherent problem in data integration and as the OBDA system operates outside the database engine, knowing all these factors is difficult or even impossible. In what follows, we propose to take this decision according to a heuristic that depends only on the size of the data and the DTR of the input table, whose estimation can be obtained using data summarization.

The main assumption that we make and will help us take decisions regarding duplicate elimination states that the impact of an input table with DTR equal to a constant number  $n$  in the number of tuples of the final query result is proportional to  $n$ . As a result of this assumption, the selectivity of the query plays the most important role regarding the duplicate elimination decisions. Intuitively, a query whose result size is much larger than the size of the intermediate result for which we examine the duplicate elimination option, it is expected to be faster if we first perform the elimination, as each tuple of the intermediate result has as impact the creation of a large number of tuples in the final result. On the other hand, when we have very selective queries with few results, whereas the size of the intermediate result under consideration is much larger, one would expect that each tuple of the intermediate result does not add that much to the total cost of the query in order to counterbalance the cost of a duplicate elimination, especially when expecting the optimizer to limit the sizes of intermediate query results as soon as possible.

**A Heuristic Regarding Duplicate Elimination.** Given a database instance  $D$ , a query  $q$  that has the form (5) and whose result over  $D$  is the relation instance  $Q$  and an input table  $I_i(\mathbf{x}_i)$  of  $q$ , then perform duplicate elimination on input table  $I_i(\mathbf{x}_i)$  prior to execution of  $q$  if

$$Size_Q - \frac{Size_Q}{DTR_{Ans}} > \frac{Size_{Ans}}{DTR_{Ans}}$$

where relation instance  $Ans$  is the result of the projection query of  $I_i(\mathbf{x}_i)$  from  $q$  on  $D$  and  $Size_Q$  and  $Size_{Ans}$  are the estimated sizes (in bytes) of relation instances  $R$  and  $A$  respectively. That is, duplicate elimination should be performed if it is expected that the reduction on the size of the final result will be bigger than the size of the intermediate result with duplicate elimination.

## 5.2 Considering Multiple Input Tables and Union Subqueries

When more than one input tables with  $DTR > 1$  exist for a query, the decision whether duplicate elimination should be performed for a specific one clearly depends on the decisions that have been taken for the rest, as the size of the result changes depending on the other decisions. In order to avoid examining all the combinations, a greedy approach is used, choosing at first the one that gives the biggest gain according to the heuristic (considering the heuristic as a fraction instead of an inequality). Then the chosen input table is excluded and a new estimation for the final result is made, considering that duplicate elimination is performed on the chosen input table and this step is repeated until none of the rest input tables gives gain according to the heuristic.

When dealing with multiple subqueries of a UNION or UNION ALL query the heuristic must be modified as follows:

$$\sum_{i=1}^n (Size_{Q_i} - \frac{Size_{Q_i}}{DTR_{Ans}}) > \frac{Size_{Ans}}{DTR_{Ans}}$$

where each  $Q_i$  is the result of a subquery  $q_i$  that uses the input table. This modified form should only be used when the the corresponding RDBMS reuses an intermediate result, for example when a temporary table is created. If the result is defined as a nested select subquery instead, then the original form of the heuristic should be used for each subquery separately.

Note that when dealing with multiple subqueries, one has to identify when the projection queries of input tables  $I_i(\mathbf{x}_i)$  coming from different subqueries are equivalent. This process is straightforward, as we are dealing with queries referencing a single input table, and as a result we can construct a single expression consisting of the relation name (for  $R_1, \dots, R_m$  in (5)) or a mapping assertion ID (for  $Aux_1, \dots, Aux_n$ ), the projected columns in ascending order and also in ascending order positions of constants in  $\mathbf{x}_i$  and the corresponding values and then perform a simple syntactic check for equality.

## 6 Duplicates Due To Redundant Self-Joins

In this section we identify a case where redundant self-joins lead to introduction of duplicates. These self-joins are redundant under standard set semantics, and could be identified by well known conjunctive query minimization methods, without taking into consideration dependencies, such as tableau minimization-known to be NP-hard ([1], Chapter 6.2). We will see that we can identify these

$$ans(f(t), g(d_2), h(a)) \leftarrow \\ movies(t, d_1), movies(t, d_2), movies(t, d_3), hasActor(t, a)$$

(a) Unfolding 1

$$ans(f(t), g(d_2), h(a)) \leftarrow movies(t, d_2), hasActor(t, a)$$

(b) Unfolding 2

Fig. 2: Unfoldings

specific cases without having to resort to general and expensive query minimization. In Section 7.1, we experimentally show how common these cases are.

Self-joins of database tables occur frequently in queries produced in SPARQL to SQL translation, as they correspond to the reconstruction of the ternary RDF statements from their reification in terms of relational tuples. Many OBDA systems eliminate this behavior when possible using semantic knowledge in the form of functional dependencies that hold in the database [22, 20]. In case no functional dependency holds for a specific database table, often an indication of a non-normalized schema, the result is the introduction of duplicate answers in the produced SQL query, as it is shown in the next example.

*Example 2.* Consider the relational tables and the mappings from Figure 1 and the SPARQL query from Figure 1d. Also consider, as in the case of Example 1, that the ontology defines the range of property *hasDirector* to be the class *Director*. This information will be used during rewriting, so that the final unfolded query has the form shown in Figure 2a. In this translation there are three instances of the base relation *movies*, each one corresponding to the translation of a SPARQL triple pattern. This query evaluated over the given database instance will return 27 answer tuples, whereas the distinct answer tuples are only 9. If the column *title* was unique in table *movies*, then these three instances would be replaced by a single one by an OBDA system that performs the mentioned optimization with respect to functional dependencies, but if it is not, or at least the database metadata do not provide this information, then the self-joins would not be avoided. The fact that self-joins are needed is already an indication of redundancy, as it implies that a non unique column (or non unique combination of columns) is used to define a class, unless the self-join is already present in the SPARQL query. One can observe that the atom *movies(t, d<sub>1</sub>)* that corresponds to the translation of the *rdf:type* triple pattern of the SPARQL query is only contributing to duplicates in the final query and it can be eliminated. The same holds for the atom *movies(t, d<sub>3</sub>)* coming from the translation of the third triple pattern. The query after the elimination is shown in Figure 2b.



An interesting observation is that under SQL bag semantics these queries produce answers with different DTR, whereas under standard CQ containment with set semantics, they are equivalent. An OBDA system that employs semantic optimization techniques for the produced query, such as query minimization using CQ containment would produce the optimized query, but to the best of our knowledge, due to the fact that such a procedure is expensive, none of the current OBDA systems choose to use that. Note that we stress that query minimization should be used for the final query, as the problem shows up after query unfolding with respect to the database mappings, as it is the case for the optimizations that use functional dependencies, so optimization techniques used during query rewriting with respect to the ontological axioms cannot eliminate this problem.

By taking advantage of the fact that we are trying to eliminate only self-joins, we can avoid the usage of a general, possibly expensive, query containment algorithm and we only need to examine specific pairs of atoms that have the same predicate, in order to find out if one of them is “covered” by the other.

Let  $CQ_1$  be the conjunctive query  $Query(\mathbf{x}) \leftarrow R_1(\mathbf{x}_1), \dots, R_n(\mathbf{x}_n)$ . Atom  $R_i(x_i^1, \dots, x_i^n)$  is *covered by* atom  $R_j(x_j^1, \dots, x_j^n)$  in  $CQ_1$  if the following three conditions hold:

- $R_i$  and  $R_j$  have the same predicate
- for each  $x_i^k$  such that  $x_i^k \in Const$ , then  $x_i^k = x_j^k$
- for each  $x_i^k$  such that  $x_i^k \in Var$  and  $x_i^k$  appears anywhere in  $\mathbf{x}, \mathbf{x}_1, \dots, \mathbf{x}_{i-1}, \mathbf{x}_{i+1}, \dots, \mathbf{x}_n$ , then  $x_i^k = x_j^k$

If  $R_i$  is covered by  $R_j$  then obviously  $CQ_1$  is equivalent to  $Query(\mathbf{x}) \leftarrow R_1(\mathbf{x}_1), \dots, R_{i-1}(\mathbf{x}_{i-1}), R_{i+1}(\mathbf{x}_{i+1}), \dots, R_n(\mathbf{x}_n)$ .

In the previous example, both atoms  $movies(t, d_1)$  and  $movies(t, d_3)$  are covered by  $movies(t, d_2)$ .

## 7 Implementation and Experimental Evaluation

We have implemented all the described optimizations in an prototype extension of Ontop version 1.18.0. Our extension is available in github<sup>4</sup>, as a fork of the official Ontop repository. All proposed optimizations were added as extra processing phases to the result of the partial datalog evaluation, before the generation of the SQL query. As Ontop uses the so called  $\mathcal{T}$ -Mappings [19] in order to emulate H-complete ABoxes and perform the tree-witness query rewriting [11] on such ABoxes, we directly use these  $\mathcal{T}$ -Mappings as our mappings  $\mathcal{M}$ . In summary,  $\mathcal{T}$ -Mappings are created from the initial mapping, by compiling in them information about ontological axioms such as property and class hierarchies.

*Experimental Setup.* All experiments were carried out on a machine with an Intel Core i7-3770K processor with 8 cores and 16 GB of RAM running UBUNTU

<sup>4</sup> <https://github.com/dbilid/ontop>

16.04. As our intention was to examine how our optimizations perform in different underlying systems, we used four different back-ends: PostgreSQL (version 9.3), MySQL (version 5.7) and two of the most widely used proprietary RDBM systems, which due to their license we will call *System I* and *System X*. All systems were setup and tuned for usage in a machine with 16GB RAM. The schema and data in all systems were identical and all the proposed indexes were created.

## 7.1 Experiments with NPD and LUBM Benchmarks

We have performed an experimental evaluation of our techniques using the LUBM [7] and NPD [13] benchmarks, with the ontology and mappings that are publicly available at the Ontop repository on github<sup>5</sup> and with existential reasoning enabled. Both datasets were generated for scale 100. In PostgreSQL database size was about 1.1 GB for LUBM and about 5.8 GB for NPD.

*Mappings.* For LUBM benchmark in total **84** mapping assertions were produced as  $\mathcal{T}$ -Mappings from Ontop. Out of these, we only needed to make DTR estimations using statistics for **15**, as we found all the others to have  $DTR=1$  according to Proposition 1. Offline gathering and building of statistics for the columns referenced in the candidate mapping assertions took **28** sec. Making DTR estimations for the result of the 15 mappings took less than **0.5** seconds. For NPD, out of **1091** produced assertions in  $\mathcal{T}$ -Mappings **112** were found to have  $DTR > 1$ . Building statistics took **45** sec. and making DTR estimations about **2** sec.

*Queries.* For LUBM we used the original 14 queries. Covered self-join elimination was applicable to **6** queries, namely: 2, 4, 7, 8, 9 and 12. Early duplicate elimination was applicable to **4** queries: 2, 4, 7 and 9. Out of these, 4 and 7 were not finally found beneficial from our heuristic. For NPD we used a subset of 19 out of the original 30 queries: queries 1-12, 22-25 and 28-30, excluding the queries that use GROUP BY, as it is not supported by the used Ontop version, queries that contain OPTIONAL and queries that produce empty SQL due to incompatible IRIs. To these queries we added four more, in order to showcase the advantage of duplicate elimination. The reason for this addition is that despite the fact that many mappings introduce duplicates, the existing queries are only using a small subset of the mappings that mostly avoid this problem. We believe that the four added queries are sensible and simple, yet their evaluation proved very hard. This showcases that the problem we are dealing with is also present in the NPD benchmark. These new queries are numbered 31 to 34 and presented in Appendix A. Out of the final 25 NPD queries, covered self-join elimination was applicable to **10** queries: 6, 9, 10, 11, 12, 23, 24, 29, 30 and 32. Early duplicate elimination was applicable to **6** queries: 9, 24, 31, 32, 33 and 34.

<sup>5</sup> <https://github.com/ontop/iswc2014-benchmark/tree/master/LUBM> and <https://github.com/ontop/npd-benchmark>

Duplicate elimination for query 9 was not performed according to our heuristic. For all the queries, pushing duplicate elimination before IRI construction was performed, according to Section 4. All SPARQL queries were executed using the DISTINCT modifier.

*Overhead in Optimization.* Total optimization time for the 14 LUBM queries total time increased from **533** ms to **948** ms, whereas for the 23 NPD queries the increase was from **748** ms to **1120** ms. The given times include the total time from parsing each SPARQL query to outputting the corresponding SQL query. The first time is the time needed by the original Ontop version, whereas the second time is the time needed by our modified version that performs all the described optimizations.

*Results.* We executed the queries with and without the proposed optimizations. For each query we used a timeout of 1000 seconds. For each setting, all queries were executed sequentially according to their numbering, after a full system reboot. The given times measure the total time needed for each query including the optimization time in Ontop, the execution time in the relational back-end and the time to obtain the results in Ontop. All the results were obtained, but they were not saved or processed otherwise. For early duplicate elimination, temporary tables were used during execution in the same session as the main query and unique indexes were created on those tables. All times are in milliseconds. For each query we present the time needed to obtain the first 1000 (that was the JDBC ResultSet fetch size used) answers (First) and all the answers (Total) without the optimizations and with the optimizations (First Opt and Total Opt). All results and the produced SQL queries are available in our repository in github<sup>6</sup>.

Results are presented in Table 1 for PostgreSQL, Table 2 from System I, Table 3 from System X and Table 4 for MySQL. For PostgreSQL without the optimizations we had three timeouts with average time (excl. timeouts) for first answers **33237** ms and for all answers **33726**. With the optimizations no timeout occurred and the times were **8005** ms for first answers and **15184** for all answers. In System I without the optimizations we had one timeout, with **73201** ms for first answers and **74135** for all answers. With the optimizations no timeout occurred and the times were **8227** ms for first answers and **14937** for all answers. System X was the only system that did not have any timeouts. Without the optimizations we had average time for first answers **35247** and for all answers **35995** ms. With the Optimizations the times were **4664** and **10789** ms. In MySQL without the optimizations we had four timeouts with average time (excl. timeouts) for first answers **31941** ms and for all answers **32322**. With the Optimizations we had two timeouts and the times (excl. timeouts) were **24197** and **24914** ms. MySQL is the only system that does not take advantage of result pipelining. Furthermore, even in the optimized setting we had two timeouts for NPD queries 32 and 34. On closer inspection, some composite indexes were

<sup>6</sup> <https://github.com/dbilid/ontop/tree/version3/results>

found to be used inefficiently. We could solve this problem by creating four new non-composite indexes, but all the times given are without this modified design.

*Considering Duplicate Elimination Only.* In order to isolate the impact of early duplicate elimination, we executed the seven queries such that this optimization is beneficial, with the optimizations for pushing IRI construction before duplicate elimination and covered atom elimination disabled. To these queries, we also added two extra NPD queries from [23] coming from user requirements, for which we also found that early duplicate elimination is applicable. These are queries USER13 and USER19. Results for all systems are presented in Table 5, where SI-O stands for System I with all optimizations enabled, SI-D for system I with only early duplicate elimination enabled and SI for the execution without any optimization in System I. Same holds for System X (SX), PostgreSQL (PG) and MySQL (MS). As before, all produced queries are available in github. From the results we can see that the impact of early duplicate elimination varies depending on each query and each system. For example, in query USER19, the decision leads to a slight increase in execution time for both systems I and X, but it is very beneficial for PostgreSQL. On a closer inspection, we saw that both these systems do take advantage of early duplicate elimination opportunities in some cases, such as when distinct rows can be obtained without extra effort using an index of a base table for all projected columns. Learning exactly when these system consider early duplicate elimination can lead to even larger improvement, if it is taken into consideration from our method as system-specific optimization. In any case, from our experiments it seems that even these systems miss many opportunities for early duplicate elimination that lead to better plans and certainly do not consider using the same duplicate elimination result for different union subqueries. As a result, even with the system-agnostic version, our method effectively eliminates most timeouts. Even in System X, where no timeout occurs in the unoptimized setting, early duplicate elimination alone leads to a decrease of **38%** in total execution time.

## 7.2 Evaluating the Duplicate Elimination Heuristic

In this section we present experimental justification for the use of our heuristic regarding duplicate elimination. For this purpose, we have chosen four query fragments from the LUBM benchmark and four from NPD, such that duplicate elimination is applicable on them, as it was found during the previously described experiments. Each query fragment consists of a single select-from-where subquery. The fragments were chosen such that they have varying characteristics regarding the execution time, the number of results and the DTR of the mapping assertion under consideration. In order to test these queries with different selectivities, we applied to them extra filters. As LUBM100 contains information about exactly 100 universities, we used a filter on the university ID attribute in direct correspondence to the percentage of selectivity, whereas for NPD we used different filters for each fragment. We used filters that result in selectivity percentage of 1, 5, 10, 30 and 60, resulting in a total of 40 queries per system. We

	First	Total	First Opt	Total Opt
L1	166	168	179	187
L2	timeout	timeout	3049	3091
L3	121	130	48	52
L4	1485	1819	191	509
L5	30	34	11	16
L6	23309	24119	1328	8855
L7	1382	1417	230	305
L8	1262	1282	193	389
L9	152640	152755	8611	47616
L10	131	134	30	35
L11	39	41	45	49
L12	7	9	5	10
L13	178	181	183	190
L14	4574	5034	14	4419
N1	5666	7445	1637	6973
N2	7002	7162	4429	4575
N3	1132	1212	986	1049
N4	33091	34272	30791	31910
N5	89	99	77	85
N6	208837	210070	36455	37754
N7	1163	1171	1147	1159
N8	570	579	400	419
N9	3298	3368	2758	2857
N10	5725	5805	5882	5985
N11	41095	41549	24818	25245
N12	89077	89952	62005	62975
N22	9132	9877	419	5075
N23	11568	12146	352	5108
N24	5603	5708	366	711
N25	15600	16813	15	15255
N28	44615	46735	9312	32549
N29	92372	92842	50108	53867
N30	170469	171448	30556	112623
N31	timeout	timeout	11725	42120
N32	190864	191260	466	513
N33	timeout	timeout	5357	43017
N34	7759	9784	1999	4248
Avg.	33237 <sup>1</sup>	33726 <sup>1</sup>	8005	15184

<sup>1</sup> excluding timeouts

Table 1: PostgreSQL Results (ms)

	First	Total	First Opt	Total Opt
L1	256	261	271	280
L2	12345	12376	8973	9038
L3	67	70	80	84
L4	17577	17875	2400	2791
L5	21	27	22	28
L6	12678	14416	1656	6703
L7	1907	1931	558	607
L8	680	705	296	389
L9	279952	280279	39605	144687
L10	382	385	34	36
L11	116	119	86	88
L12	263	268	36	40
L13	1841	1846	4377	4383
L14	2232	3175	24	2525
N1	5930	8753	4316	7210
N2	736	996	803	1064
N3	161	285	163	292
N4	2657	4907	1662	4211
N5	67	74	62	69
N6	601545	603907	17225	18755
N7	2972	2984	1150	1159
N8	220	232	153	168
N9	3585	3659	1899	2000
N10	29438	29587	4787	4924
N11	181829	182558	29676	30232
N12	427250	428831	36336	37945
N22	6688	8571	317	4142
N23	9536	11118	759	3912
N24	6785	7215	1062	1369
N25	6178	8297	12	6647
N28	33300	38518	6648	25566
N29	129934	132335	52140	55194
N30	668359	672523	63149	99597
N31	timeout	timeout	6250	29790
N32	8138	8183	3734	3761
N33	169320	170212	7508	34861
N34	10291	11396	6186	8120
Avg.	73201 <sup>1</sup>	74135 <sup>1</sup>	8227	14937

<sup>1</sup> excluding timeouts

Table 2: System I Results (ms)

	First	Total	First Opt	Total Opt
L1	202	205	100	110
L2	14214	14250	3548	3609
L3	49	51	5	11
L4	31975	32234	3529	4071
L5	31	36	16	26
L6	48010	49046	496	12370
L7	4656	4681	837	883
L8	3280	3296	235	413
L9	69468	69569	10633	32863
L10	10	12	8	12
L11	124	126	7	9
L12	23	26	8	15
L13	351	354	244	251
L14	7588	8015	15	7843
N1	13925	15879	1542	12334
N2	1046	1187	905	1047
N3	94	180	103	193
N4	8814	10045	7933	9394
N5	22	29	36	47
N6	45127	46321	26767	28103
N7	1986	1994	1906	1916
N8	400	405	215	221
N9	5650	5722	5575	5775
N10	5782	5878	7094	7230
N11	25121	25533	12970	13450
N12	48597	49412	22878	23744
N22	24595	25753	199	7653
N23	27000	27976	210	7790
N24	8015	8091	629	1173
N25	19555	20632	20	15636
N28	82159	87278	9450	34978
N29	114534	116666	18123	23699
N30	216026	220217	15327	43112
N31	169392	173868	12262	59072
N32	5146	5157	543	598
N33	294668	294786	6606	33474
N34	6532	6915	1596	6053
Avg.	35247	35995	4664	10789

Table 3: System X Results (ms)

	First	Total	First Opt	Total Opt
L1	795	797	280	289
L2	488058	488098	8431	8476
L3	270	275	211	214
L4	6869	7131	943	1248
L5	15	23	11	17
L6	30891	31117	23756	24087
L7	3347	3377	201	251
L8	13601	13612	829	845
L9	timeout	timeout	259120	259248
L10	2	5	2	6
L11	45	47	43	45
L12	3	6	2	7
L13	133	137	124	128
L14	6017	6171	282	1556
N1	14432	15272	9817	11817
N2	2339	2702	2335	2679
N3	399	541	470	612
N4	17199	19724	71046	73748
N5	37	48	40	51
N6	58664	60859	73675	75967
N7	1278	1285	1375	1390
N8	541	547	471	481
N9	10068	10156	8376	8489
N10	2980	3125	2054	2281
N11	14512	15245	49713	50483
N12	28052	29652	99437	100994
N22	15202	15545	7191	7508
N23	16627	16903	4566	6200
N24	11055	11102	2866	2987
N25	24682	25084	7513	13381
N28	51047	52097	43672	45220
N29	84224	84497	57356	59199
N30	143203	143666	102987	103497
N31	timeout	timeout	timeout	timeout
N32	timeout	timeout	5045	5068
N33	timeout	timeout	timeout	timeout
N34	7475	7664	2665	3476
Avg.	31941 <sup>1</sup>	32322 <sup>1</sup>	24197 <sup>1</sup>	24914 <sup>1</sup>

<sup>1</sup> excluding timeouts

Table 4: MySQL Results (ms)



Query	SI-O	SI-D	SI	SX-O	SX-D	SX	PG-O	PG-D	PG	MS-O	MS-D	MS
NPD24	1.37	7	8.3	1.17	7.99	8.09	0.71	5.35	5.71	2.99	10.8	11.1
NPD31	29.8	49.2	T/O	59.1	158	174	42.1	72.8	T/O	T/O	T/O	T/O
NPD32	3.76	4.31	8.18	0.6	1.46	5.16	0.51	0.72	191	5.07	8.31	T/O
NPD33	34.9	116	170	33.5	177	295	43	108	T/O	T/O	T/O	T/O
NPD34	8.12	11.2	11.4	6.05	7.44	6.92	4.25	4.23	9.78	3.48	6.23	7.66
USER13	17.7	147	157	25.7	113	315	32.7	T/O	T/O	T/O	T/O	T/O
USER19	2.38	2.85	2.16	1.54	1.88	1.44	2.79	3.16	32.3	10.1	10.9	144
LUBM2	9.04	14.1	12.4	3.61	15.4	14.3	3.09	4.27	T/O	8.48	8.28	488
LUBM9	145	214	280	32.9	65.3	69.6	47.6	78.6	153	259.25	352	T/O
Total	252	566	<649	164	547	890	177	<277	<392	<289	<397	<651

Table 5: Effect of Early Duplicate Elimination (Times in sec.)

executed each of these 40 queries with and without duplicate elimination performed, resulting in a total of 240 runs for all systems. The results were obtained with warm caches.

In the upper part of Table 6 (one-time) we present the total execution times for these queries per system, depending on the duplicate elimination strategy. The titles of the first three columns are self explanatory. The fifth column gives the total time, if always the best strategy were chosen for each system. The fourth column gives the best time, if for each query and each selectivity, the best common strategy was chosen for all systems. This way, the difference between the fourth and fifth column can give an indication of how similar the behaviors of the systems are, whereas comparison of third and fourth columns can give a measure of how well our heuristic takes advantage of this common behavior.

One can observe that the strategy of always performing duplicate elimination is much better than never performing, and that even the strategy of always choosing the best approach is not extremely better. The reason for this result is that for queries with low selectivity, the execution time is much larger and dominates the total time. For these queries, performing duplicate elimination is preferable and sometimes gives up to two orders of magnitude better results. In order to simulate a query mix such that low selectivity queries do not dominate execution time, we also computed results where we give very selective queries a weight, such that queries with 1% selectivity have been executed 60 times, queries with 5% selectivity have been executed 12 times, etc. We present the total execution time under this setting in the lower part of Table 6. Exact times and queries can be found at our github repository.

## 8 Related Work and Conclusions

In this work, we considered query reformulations such that the final SQL query is a union of CQs. In [3] a cost based comparison of different reformulations is carried out, considering that the database stores the data as triples, that is no unfolding with respect to mapping assertions is carried out. In general, the final SQL query will be a join of unions of CQs. The authors also note the problem of large intermediate results due to duplicates, and they always perform duplicate elimination on the result of each union of CQs. We believe that our work

	System	Always	Never	Heuristic	Best (common)	Best(Separate)
<i>one-time</i>	PostgreSQL	13345	168785	12854	12638	12353
	MySQL	281598	-	281685	279522	279265
	SystemI	10733	143616	9906	9693	9502
	SystemX	20558	27479	8588	8803	7280
<i>query-mix</i>	PostgreSQL	167116	618328	144984	146406	143191
	MySQL	1129311	-	1066499	1056659	1056145
	SystemI	135790	520408	102724	101984	99989
	SystemX	167761	220408	93660	90557	83045

Table 6: Query Results for Different Duplicate Elimination Strategies

can be combined with the aforementioned method, in order to push duplicate elimination to certain input tables only and avoid unnecessary operations. An extension of this work for arbitrary relational schemas is presented in [14].

In [9] the authors adopt a logic which enables them to avoid mappings when using an *object-relational* back-end and a combination of data completion and query rewriting. During this process primary keys are used for object identification, removing the need for duplicate elimination. Also, the authors use disjointness axioms in the ontology to further remove the need of duplicate elimination between unions. We believe this is orthogonal to our approach as presented in Section 4, as disjointness axioms can be used to further split query groups.

Ultrawrap-OBDA is presented in [21]. From the description of the system, it seems that the saturated mappings used will also suffer from the problem of duplicates in the case of the inference rules  $(A, \text{dom}, B)$  and  $(A, \text{range}, B)$ . After the identification of these mappings, if Ultrawrap decides to materialize them, duplicate elimination should be used, or in the case that these mappings are not materialized, we believe that our heuristic is straightforward to be applied.

[6] presents query rewriting and optimization techniques that eliminate redundant atoms during the application of a resolution based algorithm. To do so, they employ a method that takes into consideration the *tuple-generating dependencies* (TGDs) of the ontological language they consider, which unlike the DL-Lite languages, considers atoms of arbitrary arity, thus it's conceptually closer to the relational model and does not need separate mappings, so a separate unfolding phase is not needed.

We have identified duplicate answers as a bottleneck in OBDA query processing and we have proposed solutions to overcome this problem. We believe that using cost-based planning is a prominent direction towards OBDA query optimization, that has not been fully explored yet. In future work, we plan to incorporate decisions about physical database design by analyzing the mapping assertions. One more direction regarding future research has to do with duplicate elimination in case the OBDA system is equipped with query processing capabilities, in other words when it acts as a *mediator*. In this setting, along with decisions regarding which query fragments should be evaluated in external

databases, one should decide when duplicate elimination should be “pushed” to *endpoints* or performed by the OBDA processing engine during data import.

## References

1. Serge Abiteboul, Richard Hull, and Victor Vianu. *Foundations of databases: the logical level*. Addison-Wesley Longman Publishing Co., Inc., 1995.
2. Dina Bitton and David J DeWitt. Duplicate record elimination in large data files. *ACM Transactions on database systems (TODS)*, 8(2):255–265, 1983.
3. Damian Bursztyń, François Goasdoué, and Ioana Manolescu. Teaching an RDBMS about ontological constraints. *Proceedings of the VLDB Endowment*, 9(12):1161–1172, 2016.
4. Surajit Chaudhuri and Moshe Y Vardi. Optimization of real conjunctive queries. In *Proceedings of the twelfth ACM SIGACT-SIGMOD-SIGART symposium on Principles of database systems*, pages 59–70. ACM, 1993.
5. Alexandros Chortaras, Despoina Trivela, and Giorgos B Stamou. Optimized query rewriting for owl 2 ql. In *CADE*, volume 11, pages 192–206. Springer, 2011.
6. Georg Gottlob, Giorgio Orsi, and Andreas Pieris. Query rewriting and optimization for ontological databases. *ACM Transactions on Database Systems (TODS)*, 39(3):25, 2014.
7. Yuanbo Guo, Zhengxiang Pan, and Jeff Heflin. Lubm: A benchmark for owl knowledge base systems. *Web Semantics: Science, Services and Agents on the World Wide Web*, 3(2):158–182, 2005.
8. Yannis Ioannidis. The history of histograms (abridged). In *Proceedings of the 29th international conference on Very large data bases-Volume 29*, pages 19–30. VLDB Endowment, 2003.
9. Jason St Jacques, David Toman, and Grant E Weddell. Object-relational queries over  $CFD_{nc}^{\forall}$  knowledge bases: OBDA for the SQL-literate. In *Description Logics*, 2016.
10. Evgeny Kharlamov, Dag Hovland, Ernesto Jiménez-Ruiz, Davide Lanti, Hallstein Lie, Christoph Pinkel, Martin Rezk, Martin G Skjæveland, Evgenij Thorstensen, Guohui Xiao, Dmitriy Zheleznyakov, and Ian Horrocks. Ontology based access to exploration data at Statoil. In *International Semantic Web Conference*, pages 93–112. Springer, 2015.
11. Stanislav Kikot, Roman Kontchakov, and Michael Zakharyashev. Conjunctive query answering with owl 2 ql. In *Thirteenth International Conference on the Principles of Knowledge Representation and Reasoning*, 2012.
12. Henning Koehler and Sebastian Link. Armstrong axioms and boyce-codd-heath normal form under bag semantics. *Information processing letters*, 110(16):717–724, 2010.
13. Davide Lanti, Martin Rezk, Guohui Xiao, and Diego Calvanese. The NPD benchmark: Reality check for OBDA systems. In *Proc. of the 18th Int. Conf. on Extending Database Technology (EDBT)*, 2015.
14. Davide Lanti, Guohui Xiao, and Diego Calvanese. Cost-driven ontology-based data access. 2017.
15. Glenn N Paulley and Per-Åke Larson. Exploiting uniqueness in query optimization. In *Proceedings of the 1993 conference of the Centre for Advanced Studies on Collaborative research: distributed computing-Volume 2*, pages 804–822. IBM Press, 1993.

16. Jorge Pérez, Marcelo Arenas, and Claudio Gutierrez. Semantics and complexity of SPARQL. *ACM Transactions on Database Systems (TODS)*, 34(3):16, 2009.
17. Héctor Pérez-Urbina, Ian Horrocks, and Boris Motik. Efficient query answering for owl 2. *The Semantic Web-ISWC 2009*, pages 489–504, 2009.
18. Antonella Poggi, Domenico Lembo, Diego Calvanese, Giuseppe De Giacomo, Maurizio Lenzerini, and Riccardo Rosati. Linking data to ontologies. In *Journal on data semantics*, pages 133–173. Springer, 2008.
19. Mariano Rodríguez-Muro, Roman Kontchakov, and Michael Zakharyashev. Ontology-based data access: Ontop of databases. In *International Semantic Web Conference*, pages 558–573. Springer, 2013.
20. Mariano Rodríguez-Muro and Martin Rezk. Efficient SPARQL-to-SQL with R2RML mappings. *Web Semantics: Science, Services and Agents on the World Wide Web*, 33:141–169, 2015.
21. Juan F Sequeda, Marcelo Arenas, and Daniel P Miranker. OBDA: query rewriting or materialization? in practice, both! In *International Semantic Web Conference*, pages 535–551. Springer, 2014.
22. Juan F Sequeda and Daniel P Miranker. Ultrawrap: SPARQL execution on relational data. *Web Semantics: Science, Services and Agents on the World Wide Web*, 22:19–39, 2013.
23. Martin G Skjæveland, Espen H Lian, and Ian Horrocks. Publishing the norwegian petroleum directorates factpages as semantic web data. In *International Semantic Web Conference*, pages 162–177. Springer, 2013.
24. Arun Swami and K Bernhard Schiefer. On the estimation of join result sizes. In *International Conference on Extending Database Technology*, pages 287–300. Springer, 1994.

## A NPD Queries 31-34

```
SELECT DISTINCT ?q ?u
WHERE {
  ?q :inLithostratigraphicUnit ?u .
  ?u rdf:type :LithostratigraphicUnit .
}
```

Listing 1.1: Query NPD 31

```
SELECT DISTINCT ?quadrant ?name
WHERE {
  ?q rdf:type :Quadrant .
  ?quadrant :name ?name .
}
```

Listing 1.2: Query NPD 32

```
SELECT DISTINCT ?unit ?era
WHERE {
  ?unit :geochronologicEra ?era .
  ?unit rdf:type :LithostratigraphicUnit .
}
```

Listing 1.3: Query NPD 33

```
SELECT DISTINCT ?wellbore ?discovery ?year
WHERE {
  ?wellbore rdf:type :Wellbore .
  ?wellbore :wellboreForDiscovery ?discovery .
  ?discovery :discoveryYear ?year
}
```

Listing 1.4: Query NPD 34